(a) Segment-level rules: one decision per segment

Reference points: no correction · always accept · oracle

Edit-density cap: accept if the candidate changes little

Decoder confidence: mean or minimum token log-probability

Learned router: boosted trees on text, alignment and confidence features (cross-fitted on dev)

Quality estimation: document-level and trained encoder

LLM verifier: A/B choice between x and c

Learn-then-Test-style threshold on the harm rate (exploratory)

All rules score the same frozen candidates.

(b) Edit-level rules: one decision per atomic edit

OCR x `tbe qnick brown fox jumps`

candidate c `the quick brown dog jumps`

`tbe → the`	`qnick → quick`	`fox → dog`
accept	accept	reject

output `the quick brown fox jumps`

Edit scores: edit size · per-edit token log-probability · edit router · cascade (segment router, then edit router)

Edits are cut from the character alignment at unchanged whitespace.

Operating points for τ (set on dev, evaluated once on test)

CER-optimal: lowest dev CER · Harm budget: lowest dev CER while harming at most 10%, 25% or 50% as many segments as always-accept